




Deep Learning

Biological Neuron



- Massively parallel interconnected network of neurons
- Sense organs relay information to lowest layer of neurons
 - Fired neurons relay information to other connected neurons
 - Division of work – respond to certain stimulus
- Average human brain has around 10^{11} neurons


Basic Building Block



Classical Perceptron

$$\mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ x_3 \\ \dots \\ x_n \end{pmatrix}$$


$$\mathbf{w} = \begin{pmatrix} w_1 \\ w_2 \\ w_3 \\ \dots \\ w_n \end{pmatrix}$$



$$\begin{aligned} \hat{y} &= (\sum w_i x_i) \\ &= \mathbf{w}^T \mathbf{x} \end{aligned}$$

$$\forall i = 1 \dots n$$

Perceptron – Basic Building Block



Will you watch this movie???

$$\hat{y} = (\Sigma w_i x_i)$$

$$\hat{y} = 1$$
$$= 0$$

$$\forall i = 1 \dots n$$

$$\text{if } \Sigma w_i x_i \geq \theta$$

$$\text{if } \Sigma w_i x_i < \theta$$


$$\hat{y} = 1$$

$$= 0$$

$$\text{if } \Sigma w_i x_i - \theta \geq 0$$

$$\text{if } \Sigma w_i x_i - \theta < 0$$

Perceptron – Basic Building Block




$$\begin{aligned} \hat{y} &= 1 & \text{if } \sum w_i x_i - \theta &\geq 0 & \forall i = 1 \dots n \\ &= 0 & \text{if } \sum w_i x_i - \theta &< 0 \end{aligned}$$

$$\begin{aligned} \hat{y} &= 1 & \text{if } \sum w_i x_i &\geq 0 & \forall i = 0 \dots n, w_0 = -\theta, x_0 = 1 \\ &= 0 & \text{if } \sum w_i x_i &< 0 \end{aligned}$$

w_0 / b : bias, represents the prior or prejudice

Perceptron – Basic Building Block



$$\hat{y} = 1$$
$$= 0$$


$$\text{if } \sum w_i x_i \geq 0$$
$$\text{if } \sum w_i x_i < 0$$

$$\forall i = 0 \dots n, w_0 = -\theta, x_0 = 1$$

w_0 : bias

- A movie buff may have a very low threshold and may watch any movie [$\theta = 0$]
- A selective movie watcher may watch only a thriller, starring Matt Damon and directed by Nolan [$\theta = 3$]

Perceptron – Basic Building Block



$$\begin{aligned} \hat{y} &= 1 && \text{if } \sum w_i x_i \geq 0 \\ &= 0 && \text{if } \sum w_i x_i < 0 \end{aligned} \quad \forall i = 0 \dots n, w_0 = -\Theta, x_0 = 1$$


Linearly separable decision boundary: $\mathbf{w}^T \mathbf{x} = 0$

Perceptron separates input space into two, but:

- Weights (and threshold) can be learned
- Inputs are real-valued

Perceptron Learning Algorithm

```
 $P \leftarrow$  inputs with label 1;  
 $N \leftarrow$  inputs with label 0;  
Initialize  $\mathbf{w} = [w_0, w_1, w_2, \dots, w_n]^T$  randomly;  
while !convergence do  
    Pick random  $\mathbf{x} \in P \cup N$ ;  
    if  $\mathbf{x} \in P$  and  $\sum w_i * x_i < 0$  then  
         $\mathbf{w} = \mathbf{w} + \mathbf{x}$ ;  
    end  
    if  $\mathbf{x} \in N$  and  $\sum w_i * x_i \geq 0$  then  
         $\mathbf{w} = \mathbf{w} - \mathbf{x}$ ;  
    end  
end
```



Convergence- not making any more errors on training data or predictions are not changing

Perceptron Learning Algorithm

- Consider vectors \mathbf{w} and \mathbf{x}

$$\mathbf{w} = [w_0, w_1, w_2, \dots, w_n]^T$$

$$\mathbf{x} = [1, x_1, x_2, \dots, x_n]^T$$


Perceptron rule:

$$\begin{aligned} \hat{y} &= 1 && \text{if } \mathbf{w}^T \mathbf{x} \geq 0 \\ &= 0 && \text{if } \mathbf{w}^T \mathbf{x} < 0 \end{aligned}$$

- Find line $\mathbf{w}^T \mathbf{x} = 0$ which divides input space into two
- Every point \mathbf{x} on this line satisfies $\mathbf{w}^T \mathbf{x} = 0$


Perceptron Logic

- A perceptron will fire if weighted sum of inputs is greater than threshold ($b = -w_0$)




Single perceptron

- Single perceptron cannot deal with data which is not linearly separable
 - Have to be flexible with convergence statement while determining weights...
 - Till almost (say 90%) points are satisfying condition
 - Leads to few errors
 - May not be acceptable in critical real-world applications



Perceptron Logic


- A perceptron will fire if weighted sum of inputs is greater than threshold ($b = -w_0$)
 - Thresholding logic is harsh



If $x=0.51$, watch movie, if $x=0.49$, do not watch??

For real world problems, we need a smoother decision function

Sigmoid neuron



If $z = (w^T x + b) \longrightarrow \infty$, then $y = 1$


If $z = (w^T x + b) \longrightarrow -\infty$, then $y = 0$

If $z = (w^T x + b) = 0$, then $y = 0.5$

Range of sigmoid function: 0 to 1

Can be interpreted as probability



Sigmoid Neuron



$$y = \frac{1}{1 + e^{-(w^T x + b)}}$$

- Smooth
- Continuous
- Differentiable

Perceptron – Forward propagation




$$\begin{aligned}\hat{y} &= g(z) = g(\sum w_i x_i) \\ &= g(b + \sum w_i x_i) \\ &= g(b + \mathbf{W}^T \mathbf{X})\end{aligned}\quad \begin{aligned}\forall i &= 0 \dots n \\ \forall i &= 1 \dots n\end{aligned}$$

$$\text{where, } \mathbf{X} = \begin{pmatrix} x_1 \\ x_2 \\ \dots \\ x_n \end{pmatrix} \quad \mathbf{W} = \begin{pmatrix} w_1 \\ w_2 \\ \dots \\ w_n \end{pmatrix}$$


b : bias

Activation functions




$$\sigma'(z) = \sigma(z)(1 - \sigma(z))$$

(a)




$$\sigma'(z) = 1 - \sigma(z)^2$$

(b)



$$\sigma'(z) = \begin{cases} 1 & z \geq 0 \\ 0 & z < 0 \end{cases}$$

(c)




(d)


$$\sigma'(z) = \begin{cases} 1 & z \geq 0 \\ a & z < 0 \end{cases}$$

Importance - Activation functions

- Introduce non-linearity in the network
 - Allows to deal with non-linear data
 - Allows to approximate complex functions



Perceptron – Example




$$\mathbf{X} = \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} \quad \mathbf{W} = \begin{bmatrix} 2 \\ -3 \end{bmatrix} \quad b = 1, w_1 = 2, w_2 = -3$$

$$\hat{y} = g(b + \mathbf{X}^T \mathbf{W})$$

$$\hat{y} = g\left(1 + \begin{bmatrix} x_1 \\ x_2 \end{bmatrix}^T \begin{bmatrix} 2 \\ -3 \end{bmatrix}\right)$$

$$\hat{y} = g(1 + 2x_1 - 3x_2)$$




Example

- Ex. data points: (0.5, 0.2), (2.5, 0.9)
- At end of training, expect to find w^* , b^* such that
 $f(0.5) \rightarrow 0.2$, $f(2.5) \rightarrow 0.9$
- Loss = $(1/2) * \sum (y - \hat{y})^2$ given: $w = 0.5$, $b = 0$
= $(1/2) * [(0.2 - f(0.5))^2 + (0.9 - f(2.5))^2]$
= 0.073


$$y = \frac{1}{1 + e^{-(w^T x + b)}}$$

$$w = 0.5, b = 0$$




Example


- Ex. data points: $(0.5, 0.2)$, $(2.5, 0.9)$
- Hope to find a sigmoid function that the data points lie on the function




$$w = 0.5, b = 0$$



Sigmoid function



Changes w.r.t w



Changes w.r.t b

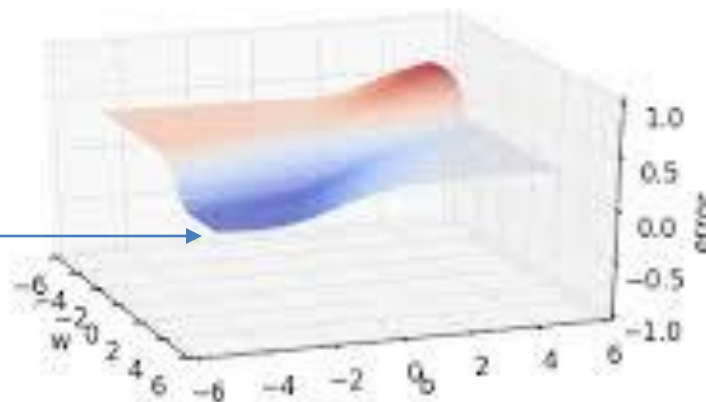
Guesses!!

w	b	Loss
0.5	0	0.073
-0.1	0	0.1481
0.94	-0.94	0.0214
1.42	-1.73	0.0028
1.65	-2.08	0.0003
1.78	-2.27	0.0000

Error surface


- 2 data points, 2 parameters (w , b)

Random search on error surface



Lowest error

Other w , b values
can be checked



- How to handle more data points and parameters??

Gradient Descent

- Find a way to traverse error surface to reach minimum value quickly
- Parameters: $\theta = [w, b]$
- Change in parameters: $\Delta\theta = [\Delta w, \Delta b]$
- New parameters: $\theta_{\text{new}} = \theta + \Delta\theta = [w_{\text{new}}, b_{\text{new}}]$
 - Change with a small stride η
$$w_{\text{new}} = w + \eta \cdot \Delta w$$
$$b_{\text{new}} = b + \eta \cdot \Delta b$$
- How to choose Δw and Δb ?

Gradient Descent

- According to Taylor series:

$$\begin{aligned} L(\theta + \eta\Delta\theta) &= L(\theta) + \eta^*(\Delta\theta)^\top \frac{\partial L(\theta)}{\partial \theta} + (\eta^2/2!)*(\Delta\theta)^\top \frac{\partial^2 L(\theta)}{\partial \theta^2}(\Delta\theta) + \dots \\ &= L(\theta) + \eta(\Delta\theta)^\top \frac{\partial L(\theta)}{\partial \theta} \quad (\text{if } \eta \text{ is small}) \end{aligned}$$

- New loss should be less than old loss
- $(\eta\Delta\theta)$ would be favourable only if:

$$L(\theta + \eta\Delta\theta) < L(\theta)$$

$$L(\theta + \eta\Delta\theta) - L(\theta) < 0$$

- Implies: $(\Delta\theta)^\top \frac{\partial L(\theta)}{\partial \theta} < 0$ (η is positive constant)


Gradient Descent

- Desired: $(\Delta\theta)^\top \frac{\partial L(\theta)}{\partial \theta} < 0$
 - Want $\Delta\theta$ to be as negative as possible
 - Dot product of two vectors: product of their magnitudes multiplied by the cosine of the angle between them
- Let θ be the angle between $\Delta\theta$ and $\frac{\partial L(\theta)}{\partial \theta}$

$$\text{Then: } -1 \leq \cos(\theta) = \frac{(\Delta\theta)^\top \frac{\partial L(\theta)}{\partial \theta}}{||\Delta\theta|| * ||\frac{\partial L(\theta)}{\partial \theta}||} \leq 1$$

Most negative: $\cos(\theta) = -1$ when $\theta = 180^\circ$

-> $\Delta\theta$ should be such that it is at 180° to the gradient $\frac{\partial L(\theta)}{\partial \theta}$



$$\cos \theta = \frac{a \cdot b}{|a| |b|}$$

Gradient Descent

- $(\Delta\theta)^\top \frac{\partial L(\theta)}{\partial \theta} < 0$
- Move in direction opposite to gradient (180° w.r.t. the gradient)

$$w_{t+1} = w_t - \eta \frac{\partial L(w,b)}{\partial w}$$

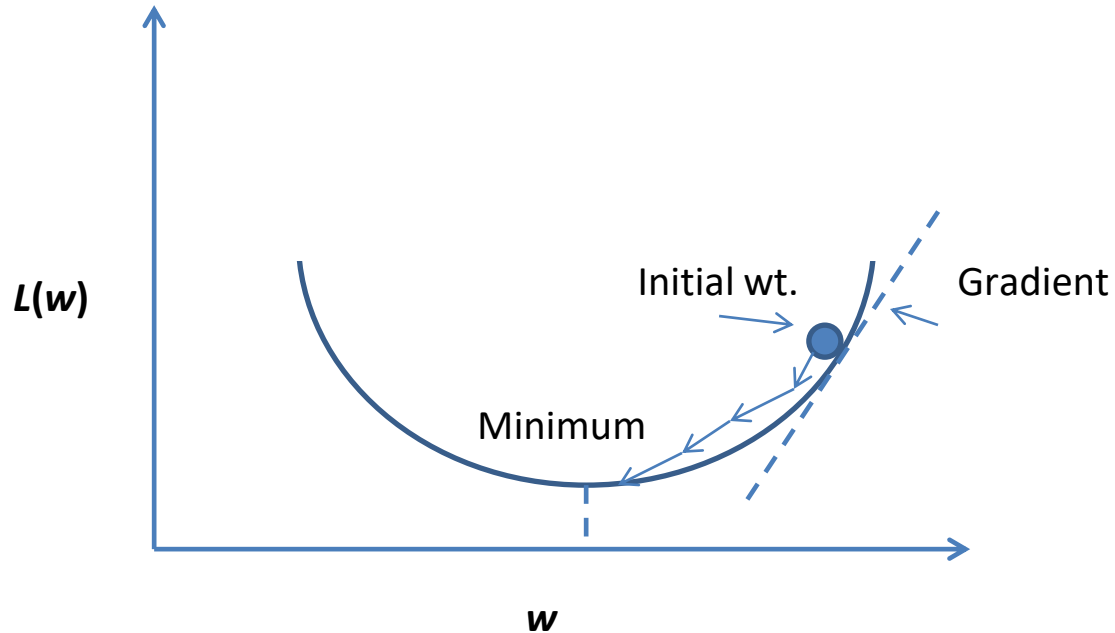
at $w = w_t$ and $b = b_t$

$$b_{t+1} = b_t - \eta \frac{\partial L(w,b)}{\partial b}$$

at $w = w_t$ and $b = b_t$

- Repeat till convergence

Gradient Descent



$$w := w - \eta * \partial L(w) / \partial w$$

Randomly pick (w_0, w_1)

Compute L

Compute gradient
 $\partial L(w) / \partial w$ - gives ascent
at that point

Take small step in
opposite direction of
gradient

Repeat until
convergence

Gradient Descent

Assuming, $L(w, b) = (1/2) * \sum (\hat{y} - y)^2$

$$\hat{y} = f(x) = 1/(1 + e^{-(wx+b)})$$

Assuming one data point only:

$$\begin{aligned}\frac{\partial L(w, b)}{\partial w} &= (1/2) * [2 * (\hat{y} - y) * \frac{\partial (\hat{y} - y)}{\partial w}] \\ &= (\hat{y} - y) * \frac{\partial}{\partial w} \left[\frac{1}{1 + e^{-(wx+b)}} \right]\end{aligned}$$

$$\frac{\partial L(w, b)}{\partial w} = (\hat{y} - y) * \hat{y} * (1 - \hat{y}) * x$$

$$\begin{aligned}\frac{\partial L(w, b)}{\partial b} &= (\hat{y} - y) * \frac{\partial}{\partial b} \left[\frac{1}{1 + e^{-(wx+b)}} \right] \\ &= (\hat{y} - y) * \hat{y} * (1 - \hat{y})\end{aligned}$$

$$\partial/\partial w [1/(1 + e^{-(wx+b)})] = f(x) * (1 - f(x)) * x$$

Gradient Descent

$$\frac{\partial L(w,b)}{\partial w} = \sum [(\hat{y}_i - y_i) * \hat{y}_i * (1 - \hat{y}_i) * x_i] \quad \text{for all points}$$

$$\frac{\partial L(w,b)}{\partial b} = \sum [(\hat{y}_i - y_i) * \hat{y}_i * (1 - \hat{y}_i)] \quad \text{for all points}$$

- Algorithm:

1. Initialize weights randomly


2. Loop until convergence:

1. Compute gradient $\frac{\partial L(w,b)}{\partial w}, \frac{\partial L(w,b)}{\partial b}$

2. Update weights $\mathbf{w} := \mathbf{w} - \eta * \frac{\partial L(w,b)}{\partial w}$
 $\mathbf{b} := \mathbf{b} - \eta * \frac{\partial L(w,b)}{\partial b}$

Return weights

Gradient Descent




Global minimum

Assuming $J(W)$ to be a convex function

Representative Power of Multilayer Networks

- A multilayer network of **sigmoid neurons** with a single hidden layer can be used to approximate any continuous function to any desired precision

Single Hidden Layer Neural Network




Hidden Layer: States of nodes are unobserved


Inputs are densely connected to perceptrons, hence they are called **Dense** layers or **Fully Connected** layers

Feedforward Neural Network

- Input is an n -dimensional vector (0^{th} layer) $\in \mathbb{R}^n$
- Network has $L-1$ hidden layers
- 1 output layer containing k neurons (ex. for k classes)
- Each neuron – aggregation and activation




Feedforward Neural Network



Assuming n^i neurons in hidden layer h^i , $sizeof(W^i) = n^i * n^{i-1}$ and $sizeof(b^i) = n^i$ between layers $i-1$ and i for $0 < i < L$


$sizeof(W^L) = n^k * n^{i-1}$ and $sizeof(b^L) = n^k$ between last hidden layer and output layer



Aggregation at layer i : $\mathbf{z}^i = \mathbf{W}^i \mathbf{a}^{i-1} + \mathbf{b}^i$

For first hidden layer: $\mathbf{z}^1 = \mathbf{W}^1 \mathbf{a}^0 + \mathbf{b}^1$

$$\begin{pmatrix} z_1^1 \\ z_2^1 \\ \dots \\ z_5^1 \end{pmatrix} = \begin{pmatrix} W_{11} & W_{12} & W_{13} \\ W_{21} & W_{22} & W_{23} \\ \dots & \dots & \dots \\ W_{51} & W_{52} & W_{53} \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} + \begin{pmatrix} b_1 \\ b_2 \\ \dots \\ b_5 \end{pmatrix} = \begin{pmatrix} \sum W_{1i} x_i + b_1 \\ \sum W_{2i} x_i + b_2 \\ \dots \\ \sum W_{5i} x_i + b_5 \end{pmatrix}$$




Activation at layer $i = g(z^i) = g(b^i + W^i a^{i-1})$

For first hidden layer: $g(z^1) = g(b^1 + W^1 a^0)$

$$\begin{pmatrix} a_1 \\ a_2 \\ \dots \\ a_5 \end{pmatrix} = \begin{pmatrix} g(z_1) \\ g(z_2) \\ \dots \\ g(z_5) \end{pmatrix}$$


Eg. $g(z_1) = \sigma(z_1) = 1 / (1 + e^{-z_1})$

g: activation function (logistic, tanh, linear etc.)




For second hidden layer: $\mathbf{z}^2 = \mathbf{W}^2 \mathbf{a}^1 + \mathbf{b}^2$

$$\begin{pmatrix} z_1^2 \\ z_2^2 \\ z_3^2 \\ z_4^2 \end{pmatrix} = \begin{pmatrix} W_{11} & W_{12} & W_{13} & W_{14} & W_{15} \\ W_{21} & W_{22} & W_{23} & W_{24} & W_{25} \\ W_{31} & W_{32} & W_{33} & W_{34} & W_{35} \\ W_{41} & W_{42} & W_{43} & W_{44} & W_{55} \end{pmatrix} \begin{pmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \\ a_5 \end{pmatrix} + \begin{pmatrix} b_1 \\ b_2 \\ b_3 \\ b_4 \end{pmatrix} = \begin{pmatrix} \sum W_{1i} a_i + b_1 \\ \sum W_{2i} a_i + b_2 \\ \sum W_{3i} a_i + b_3 \\ \sum W_{4i} a_i + b_4 \end{pmatrix}$$



Activation at layer 2 = $g(z^2) = g(b^2 + W^2 a^1)$

$$\begin{bmatrix} a_1 \\ a_2 \\ a_3 \\ a_4 \end{bmatrix} = \begin{bmatrix} g(z_1) \\ g(z_2) \\ g(z_3) \\ g(z_4) \end{bmatrix}$$



Aggregation at output layer $L = z^L = \mathbf{W}^L \mathbf{a}^{L-1} + \mathbf{b}^L$

$$z_1 = w_{11}a_1 + w_{12}a_2 + w_{13}a_3 + w_{14}a_4 + b$$


$$z_2 = w_{21}a_1 + w_{22}a_2 + w_{23}a_3 + w_{24}a_4 + b$$

$$z_3 = w_{31}a_1 + w_{32}a_2 + w_{33}a_3 + w_{34}a_4 + b$$

Activation at output layer $L = \hat{\mathbf{y}} = g(z^L) = g(\mathbf{W}^L \mathbf{a}^{L-1} + \mathbf{b}^L)$

$$\begin{bmatrix} \hat{y}_1 \\ \hat{y}_2 \\ \hat{y}_3 \end{bmatrix} = \begin{bmatrix} g(z_1) \\ g(z_2) \\ g(z_3) \end{bmatrix}$$

Neural Network Function



$$f^{(L)} (f^{(L-1)} (.... f^{(i)} (f^{(2)} (f^{(1)} (x))))))$$

Learning parameters

In given example, dimensions of parameters:

- $\mathbf{W}^1: n^1 * n$ $\mathbf{b}^1: n^1$
- $\mathbf{W}^2: n^2 * n^1$ $\mathbf{b}^2: n^2$
- $\mathbf{W}^L: k * n^2$ $\mathbf{b}^L: k$
- Assuming L layers and n^i neurons in hidden layer h^i and k neurons in output layer, no. of parameters to be learned:
 - Weights: $(L-1) * (n^i * n^{i-1}) + (n^2 * k)$ for $0 < i < L$
 - Bias: $(L-1) * n^i + k$

Learning parameters

- **Data:** $\{x_i, y_i\}$ $i = 1..m$

- **Model:**

$$\hat{\mathbf{y}} = f(\mathbf{x}) = g(\mathbf{W}^3 g(\mathbf{W}^2 g(\mathbf{W}^1 \mathbf{x} + \mathbf{b}^1) + \mathbf{b}^2) + \mathbf{b}^3)$$

$$\hat{\mathbf{y}} = [\hat{y}^1 \quad \hat{y}^2 \quad \dots \quad \hat{y}^k]$$

- **Algorithm:** Gradient Descent with back Propagation
- **Loss/Error function:** Sum of squared error loss

$$\min \frac{1}{N} \sum_{i=1}^m \sum_{j=1}^k (\hat{y}_j^i - y_j^i) \quad \text{for } i^{th} \text{ sample for all classes } j$$

Learning parameters

- Gradient Descent:

$t:=0$;

$max_iterations:=1000$;

Initialize $\theta_0 := [\mathbf{W}_0^1, \dots, \mathbf{W}_0^L, \mathbf{b}_0^1 \dots \mathbf{b}_0^L]$;

while $t++ < max_iterations$ do

$\theta_{t+1} := \theta_t - \eta \nabla \theta_t$;

end

where, $\nabla \theta_t = \left[\frac{\partial L(\theta)}{\partial W_t}, \frac{\partial L(\theta)}{\partial b_t} \right]^T$

$\nabla \theta$ composed of:

– $\nabla W^1, \nabla W^2, \dots, \nabla W^{L-1} \in \mathbb{R}^{n(i-1) \times n_i}$, $\nabla W^L \in \mathbb{R}^{n \times k}$

– $\nabla b^1, \nabla b^2, \dots, \nabla b^{L-1} \in \mathbb{R}^{n_i}$, $\nabla b^L \in \mathbb{R}^k$

Loss function

- Loss function should capture how much \hat{y}_i deviates from y_i
- $y_i \in \mathbb{R}^n$ then squared error loss can be used:

$$L(\theta) = (1/2m) * \sum (y_i - \hat{y}_i)^2$$

- Problems with squared error loss:

$$\frac{\partial L(w, b)}{\partial w} = (\hat{y} - y) * \hat{y} * (1 - \hat{y}) * x$$

- If $y_i = 1$ and $\hat{y}_i \sim 0$, $\frac{\partial L(w, b)}{\partial w} \sim 0$ Undesirable
- If $y_i = 0$ and $\hat{y}_i \sim 1$, $\frac{\partial L(w, b)}{\partial w} \sim 0$ Undesirable
- Weight updation becomes very slow

Loss function

- Cross-entropy: gives a measure on how close a predicted distribution is to a true distribution
 - True distribution p_i , Estimated distribution q_i
 - Estimated information content = $-\sum p_i \log_e(q_i)$
 - Capture difference between two probability distributions
 - If prediction is close to actual, cross entropy will be low

$$L(\theta) = -\sum y_c \log_e(\hat{y}_c)$$

$$y_c = 1$$

$$= 0$$

for all k classes

if $c = t$ (true class)

otherwise

$$L(\theta) = -\log_e(\hat{y}_t)$$

Loss function

- Objective function for classification:
 - Cross-entropy Loss

minimize: $L(\theta) = -\log_e(\hat{y}_t)$

\hat{y}_t : predicted probability of correct event

$\log_e(\hat{y}_t)$: probability that x belongs to t^{th} class, log-likelihood of data


Output Activation Function

- Binary classification:
 - Single neuron in output layer
 - Sigmoid activation function
 - Output between 0-1
 - Above threshold → One class
 - Below threshold → Another class

Output Activation Function

- Output activation function for multi-classification:
 - Sum of outputs should be 1
 - \hat{y} should be a probability distribution
 - Sigmoid – probabilities will be $0 < p < 1$ but sum not equal to 1

$$y_i = \{1 \quad 0 \quad 0 \quad 0\}$$



Classification problem

Output Activation Function

- Softmax function

$$z^L = b^L + W^L a^{L-1}$$

$$\hat{y} = g(z_j^L) = e^{z_j} / \sum e^{z_j} \quad \text{for } j = 1..k$$

z_j^L is j^{th} element of z^L

- Example: $z^L = [10 \quad 20 \quad -30]$

$$\hat{y} = [e^{10}/(e^{10} + e^{20} + e^{-30}) \quad e^{20}/(e^{10} + e^{20} + e^{-30}) \quad e^{-30}/(e^{10} + e^{20} + e^{-30})]$$

NOTE: Exponent converts –ve values to +ve values

Loss function


	Outputs	
	Real values	Probabilities
Output activation	Linear	Softmax
Loss function	Squared error	Cross-entropy

Backpropagation

How to compute $\nabla\theta$ composed of:

$$\nabla W^1, \nabla W^2, \dots, \nabla W^{L-1} \in \mathbb{R}^{n \times n}, \nabla W^L \in \mathbb{R}^{n \times k}$$


$$\nabla b^1, \nabla b^2, \dots, \nabla b^{L-1} \in \mathbb{R}^n, \nabla b^L \in \mathbb{R}^k$$



Backpropagation


Assuming classification problem, $L(\theta) = -\log_2(\hat{y}_t)$

- To learn weight w_{12}^1 use SGD and compute $\frac{\partial L(w,b)}{\partial w_{12}^1}$



Backpropagation

Assume a deep thin network, who is responsible for the loss??




Find derivative by chain rule:

$$\frac{\partial L(\theta)}{\partial W_{11}^1} = \underbrace{\frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^L}}_{\text{Output layer}} * \underbrace{\frac{\partial z_1^L}{\partial a_1^2} * \frac{\partial a_1^2}{\partial z_1^2}}_{\text{Previous hidden layer}} * \underbrace{\frac{\partial z_1^2}{\partial a_1^1} * \frac{\partial a_1^1}{\partial z_1^1}}_{\text{Previous hidden layer}} * \underbrace{\frac{\partial z_1^1}{\partial W_{11}^1}}_{\text{Weights}}$$

If we change W_{11} , how much does the loss change

Backpropagation

Assume a deep thin network




Find derivative by chain rule:

$$\begin{aligned} \frac{\partial L(\theta)}{\partial W_{11}^1} &= \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^L} * \frac{\partial z_1^L}{\partial a_1^2} * \frac{\partial a_1^2}{\partial z_1^2} * \frac{\partial z_1^2}{\partial a_1^1} * \frac{\partial a_1^1}{\partial z_1^1} * \frac{\partial z_1^1}{\partial W_{11}^1} \\ \frac{\partial L(\theta)}{\partial W_{11}^2} &= \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^L} * \frac{\partial z_1^L}{\partial a_1^2} * \frac{\partial a_1^2}{\partial z_1^2} * \frac{\partial z_1^2}{\partial W_{11}^2} \\ \frac{\partial L(\theta)}{\partial W_{11}^L} &= \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^L} * \frac{\partial z_1^L}{\partial W_{11}^L} \end{aligned}$$

BACKPROPAGATION WITH SIGMOID OUTPUT ACTIVATION & BINARY CROSS-ENTROPY LOSS

Backpropagation




Assuming binary cross-entropy function

$$L = -y \log \hat{y} - (1 - y) \log (1 - \hat{y})$$

$$\frac{\partial L(\theta)}{\partial \hat{y}} = \frac{-y}{\hat{y}} + \frac{1 - y}{1 - \hat{y}}$$

$$\frac{\partial L(\theta)}{\partial z} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z} = \left(\frac{-y}{\hat{y}} + \frac{1 - y}{1 - \hat{y}} \right) * \hat{y}(1 - \hat{y}) = \hat{y} - y$$

$$\frac{\partial L(\theta)}{\partial w} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z} * \frac{\partial \hat{z}}{\partial w} = (\hat{y} - y) * x$$



$$\frac{\partial L(\theta)}{\partial \hat{y}} = \frac{-y}{\hat{y}} + \frac{1-y}{1-\hat{y}}$$

$$\frac{\partial L(\theta)}{\partial z^2} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z^2} = \left(\frac{-y}{\hat{y}} + \frac{1-y}{1-\hat{y}} \right) * \hat{y}(1 - \hat{y}) = \hat{y} - y \quad :\delta^L = \frac{\partial L}{\partial \hat{y}} * \sigma'(z^L)$$

$$\frac{\partial L(\theta)}{\partial w^2} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z^2} * \frac{\partial z^2}{\partial w^2} = (\hat{y} - y) * a^1 \quad :\delta^L * a^1$$


$$\frac{\partial L(\theta)}{\partial a^1} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z^2} * \frac{\partial z^2}{\partial a^1} = (\hat{y} - y) * w^2 \quad : \delta^L * w^2$$

$$\frac{\partial L(\theta)}{\partial z^1} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z^2} * \frac{\partial z^2}{\partial a^1} * \frac{\partial a^1}{\partial z^1} = (\hat{y} - y) * w^2 * a^1(1 - a^1) \quad :\delta^1 = \delta^L w^2 * \sigma'(z^1)$$

$$\frac{\partial L(\theta)}{\partial w^1} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z^2} * \frac{\partial z^2}{\partial a^1} * \frac{\partial a^1}{\partial z^1} * \frac{\partial z^1}{\partial w^1} = (\hat{y} - y) * w^2 * a^1(1 - a^1) * x \quad :\delta^1 * x$$

Backpropagation Equations

- $\delta^L = \frac{\partial L}{\partial \hat{y}} \odot \sigma'(z^L)$
- $\delta^l = ((w^{l+1})^T \delta^{l+1}) \odot \sigma'(z^l)$
- $\frac{\partial L}{\partial b_j^l} = \delta_j^l$
- $\frac{\partial L}{\partial w_{jk}^l} = a_k^{l-1} \delta_j^l$



Assuming:
Initial $w_{11}^2 = 12, w_{11}^1 = 0.1$
 $\eta = 0.01$

During forward propagation:
 $a_1^1 = 0.37, z_1^1 = 0.5$


$$\frac{\partial L}{\partial w_{11}^2} = \frac{\partial L}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^2} * \frac{\partial z_1^2}{\partial w_{11}^2}$$

$$\frac{\partial L(\theta)}{\partial \hat{y}} = \frac{-y}{\hat{y}} + \frac{1-y}{1-\hat{y}} = \frac{-1}{0.25} = -4$$

$$\frac{\partial L(\theta)}{\partial z_1^2} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^2} = \hat{y} - y = 0.25 - 1 = -0.75$$

$$\frac{\partial L(\theta)}{\partial w_{11}^2} = \frac{\partial L(\theta)}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^2} * \frac{\partial z_1^2}{\partial w_{11}^2} = -0.75 * a_1^1 = -0.75 * 0.37 = -0.2775$$

$$w_{11}^{2*} = w_{11}^2 - \eta * \frac{\partial L}{\partial w_{11}^2} = 12 - 0.01 * (-0.2775) = 12.0028$$



$$\frac{\partial L}{\partial w_{11}^1} = \frac{\partial L}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial z_1^2} * \frac{\partial z_1^2}{\partial a_1^1} * \frac{\partial a_1^1}{\partial z_1^1} * \frac{\partial z_1^1}{\partial w_{11}^1}$$

$$= -0.75 * w_{11}^2 * a_1^1 (1 - a_1^1) * x_1$$

$$= -0.75 * 12 * 0.37 (1 - 0.37) * 60$$

$$= -125.874$$

$$w_{11}^{1*} = w_{11}^1 - \eta * \frac{\partial L}{\partial w_{11}^1} = 0.1 - 0.01 * (-125.874) = 1.35$$

Assignment: Compute $\frac{\partial L}{\partial w_{13}^1}$