

# Encoder-Decoder


Autoencoders for Computer Vision

# Autoencoders (AE)

- A powerful tool used in machine learning for:
  - Feature extraction
  - Data compression
  - Image reconstruction
- Used for unsupervised learning tasks
- An AE model has the ability to automatically learn complex features from input data
  - Popular method for improving accuracy of classification and prediction tasks

# Autoencoders

- Autoencoders are neural networks
  - Can learn to compress and reconstruct input data, such as images, using a hidden layer of neurons
  - Learn data encodings in an **unsupervised** manner
- Consists of two parts:
  - Encoder: takes input data and compresses it into a lower-dimensional representation called **latent space**
  - Decoder: reconstructs input data from latent space representation
- In an optimal scenario, autoencoder performs as close to perfect reconstruction as possible




# Loss Function and Reconstruction Loss

- Loss functions - critical role in training autoencoders and determining their performance
  - Most commonly used is **reconstruction loss**
  - Used to measure difference between model input and output
  - **Reconstruction error** calculated using various loss functions such as mean squared error, binary cross-entropy, or categorical cross-entropy
  - Applied method depends on type of data being reconstructed
- Reconstruction loss used to update weights of network during backpropagation to minimize difference between input and output
  - Goal: achieve low reconstruction loss
  - Low loss → model can effectively capture salient features of input data and reconstruct it accurately

# AE in Computer Vision

- Input is an image and output is a reconstructed image
  - Input image typically represented as a matrix of pixel values
  - Can be of any size, but is typically normalized to improve performance



# Encoder

- **Encoder**: Compresses input image into a lower-dimensional representation, known as **latent space** ("bottleneck" or "code")
- Encoder is:
  - Series of convolutional layers
  - Followed by pooling modules or simple linear layers, that extract different levels of features from input image
- Each layer:
  - Applies a set of filters to input image
  - Outputs a feature map that highlights specific patterns and structures in image

# Encoder

- Every convolution layer composed of  $n$  convolution filters,  $\{F_1^{(1)} \dots F_n^{(1)}\}$  each with depth  $D$
- Convolution of input volume  $I = \{I_1, \dots, I_D\}$  with filters produces  $n$  activation maps

$$O_m(i, j) = a\left(\sum_{d=1}^D \sum_{u=-2k-1}^{2k+1} \sum_{v=-2k-1}^{2k+1} F_{m_d}^{(1)}(u, v) I_d(i - u, j - v)\right) \quad m = 1, \dots, n$$

- $O(i, j)$  is pixel at position  $(i, j)$
- $2k+1$  is side of a square, odd convolutional filter
- Every convolution wrapped by a non-linear function  $a$



# Encoder

$$z_m = O_m = a(I * F_m^{(1)} + b_m^{(1)}) \quad b_m \text{ is bias for } m^{\text{th}} \text{ feature map}$$

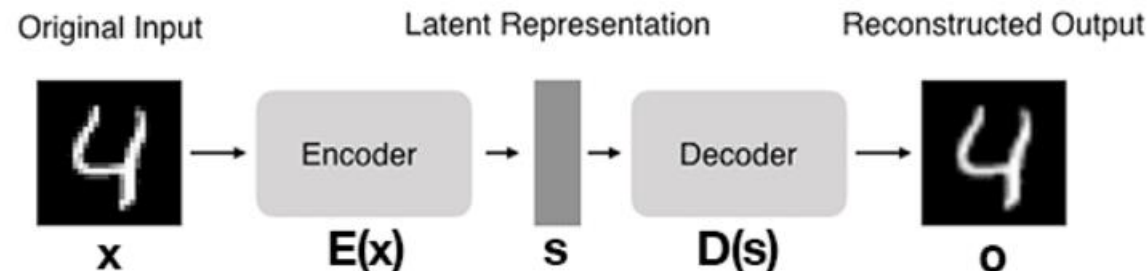
- Produced activation maps are encoding of input  $I$  in a low-dimensional space
- Convolution reduces output's spatial extent
  - Not possible to reconstruct volume with same spatial extent as input
  - **Input padding** such that  $\dim(I) = \dim(\text{decode}(\text{encode}(I)))$

# Bottleneck

- **Bottleneck / Latent Representation:** Output of encoder is a compressed representation of input image in latent space
  - Captures most important features of input image
  - Typically a smaller dimensional representation of input image
  - Restricts flow of information to decoder from encoder - allowing only most vital information to pass through
  - Prevents neural network from memorizing input and overfitting data
  - Smaller the code, lower the risk of overfitting
  - Layer generally implemented as a linear layer or as a tensor if we use convolutions
  - If input data denoted as  $\mathbf{x}$ , then latent space representation  $\mathbf{s} = E(\mathbf{x})$

# Decoder

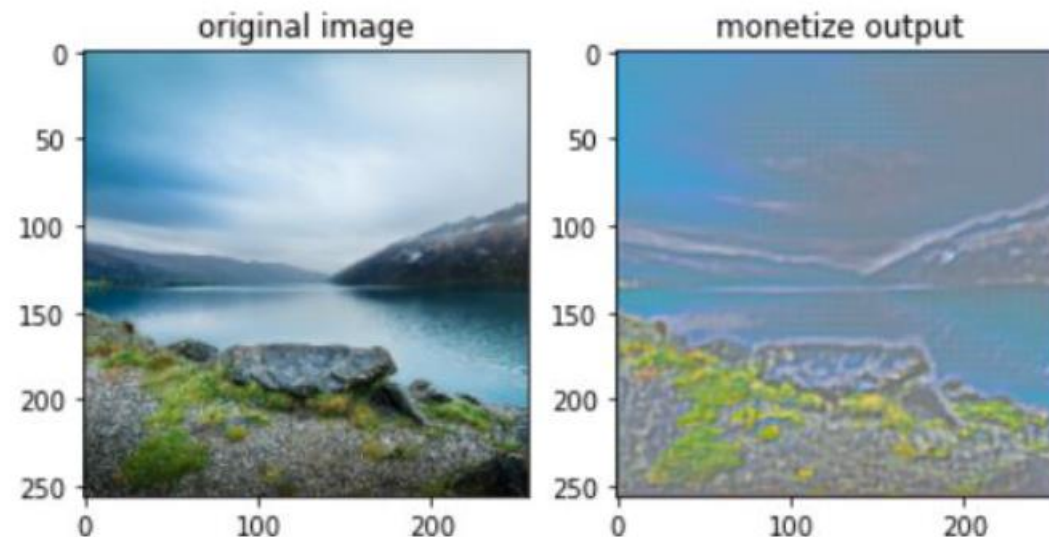
- **Decoder**: reconstructs input image from latent representation
  - Decoder is a set of multiple deconvolutional layers that gradually increase size of feature maps until final output is same size as input image
    - Usually implemented with transposed convolutions if we work with images or linear layers
  - Every layer applies a set of filters that up-sample feature maps
  - Output compared with ground truth
  - If output of decoder is  $\mathbf{o}$ , then  $\mathbf{o} = D(\mathbf{s}) = D(E(\mathbf{x}))$



# Decoder

- Output of decoder is a reconstructed image similar to input image
  - Reconstructed image may not be identical to input image
    - Since AE learns to capture most important features of input image in latent representation
  - These features can be used for tasks such as image classification, object detection, and image retrieval

AE converting input to a Monet style painting



<https://www.analyticsvidhya.com/blog/2021/01/auto-encoders-for-computer-vision-an-endless-world-of-possibilities/>

# Decoder

- $n$  feature maps  $z_{m=1,..n}$  (latent representations) produced from Encoder used as input to decoder to reconstruct image  $I$
- Hyper-parameters of decoding convolution fixed by encoding architecture
  - Filters volume  $F^{(2)}$  with dimensions  $(2k + 1, 2k + 1, n)$  to produce same spatial extent of  $I$
  - Number of filters to learn:  $D$
- Reconstructed image  $\tilde{I}$  result of convolution between feature maps  $Z = \{z_{i=1}\}^n$  and  $F^{(2)}$ 
$$\tilde{I} = a(Z * F_m^{(2)} + b^{(2)})$$
- Loss function  $L(I, \tilde{I}) = \frac{1}{2} ||I - \tilde{I}||_2^2$

# Dimensionality Reduction

- Dimensionality reduction - process of reducing number of dimensions in encoded representation of input data
- AE can learn to perform dimensionality reduction:
  - Training encoder to map input data to a lower-dimensional latent space
  - Decoder trained to reconstruct original input data from latent space representation
  - Size of latent space typically much smaller than size of input data - allowing for efficient storage and computation of data
- Through dimensionality reduction, AE can also help to remove noise and irrelevant features
  - Useful for improving performance of downstream tasks such as data classification or clustering

# Contractive autoencoder

- Designed to learn a compressed representation of input data while being resistant to small perturbations in input
  - Achieved by adding a regularization term to training objective
  - This term penalizes network for changing output with respect to small changes in input
$$Loss = L(I, \tilde{I}) + \text{regularizer}$$
- Can add a scaling parameter in front of regularization term to adjust trade-off between the two objectives

# Undercomplete Autoencoder


- Constrain number of nodes present in hidden layer(s) of network
  - Limit amount of information that can flow through the network
  - Model can learn most important attributes of input data and how to best reconstruct original input from an "encoded" state
  - For higher dimensional data, autoencoders capable of learning a complex representation of data (manifold)
  - Can be used to describe observations in a lower dimensionality and correspondingly decoded into original input space



Input layer


Hidden layers

Output layer



# Sparse Autoencoder

- Added constraint on encoding process
  - Encoder network trained to produce sparse encoding vectors - have many zero values
    - Does not require reduction in number of nodes at hidden layer
- Forces network to identify only most important features of input data
  - Construct loss function to penalize *activations* within a layer
  - Encourage network to learn an encoding and decoding which only relies on activating a small number of neurons
  - Different approach towards regularization - we normally regularize *weights* of a network, not activations
- Sensitize individual hidden layer nodes toward specific attributes of input
  - Forced to selectively activate regions of network depending on input data



- Opacity of a node corresponds with level of activation
- Individual nodes of a trained model which activate are *data-dependent*
- Different inputs will result in activations of different nodes through network

# Sparse Autoencoder

- Two main ways to impose sparsity constraint
  - Both involve measuring hidden layer activations for each training batch and adding some term to loss function to penalize excessive activations
- **L1 Regularization**: Add term to loss function to penalize absolute value of vector of activations  $a$  in layer  $h$  for observation  $i$ , scaled by tuning parameter  $\lambda$

$$L(I, \hat{I}) + \lambda \sum_i |a_i^{(h)}|$$

# Sparse Autoencoder

- **KL-Divergence**: measure of difference between two probability distributions
  - Define a sparsity parameter  $\rho$  which denotes average activation of a neuron over a collection of samples

- This expectation can be calculated as

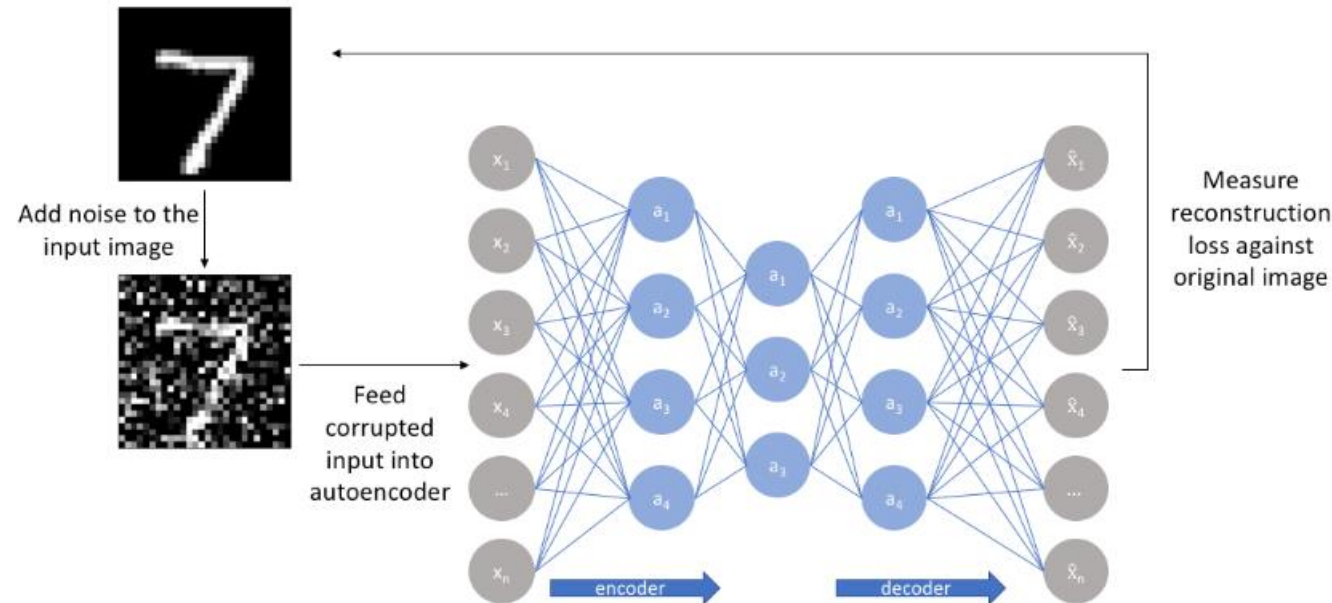
$$\hat{\rho}_j = \frac{1}{m} \sum_i [a_i^{(h)}(x)]$$

- For  $j^{\text{th}}$  neuron in layer  $h$ , activations for  $m$  training observations denoted individually as  $x$
- By constraining average activation of a neuron over a collection of samples - encouraging neurons to only fire for a subset of observations
- compare ideal distribution  $\rho$  to observed distributions  $\hat{\rho}$  over all hidden layer nodes

$$L(I, \hat{I}) + \sum_j KL(\rho || \hat{\rho})$$

# Denoising autoencoder

- Designed to learn to reconstruct an input from a corrupted version of input
  - Corrupted input created by adding noise to original input
  - Network trained to remove noise and reconstruct original input



# Denoising autoencoder

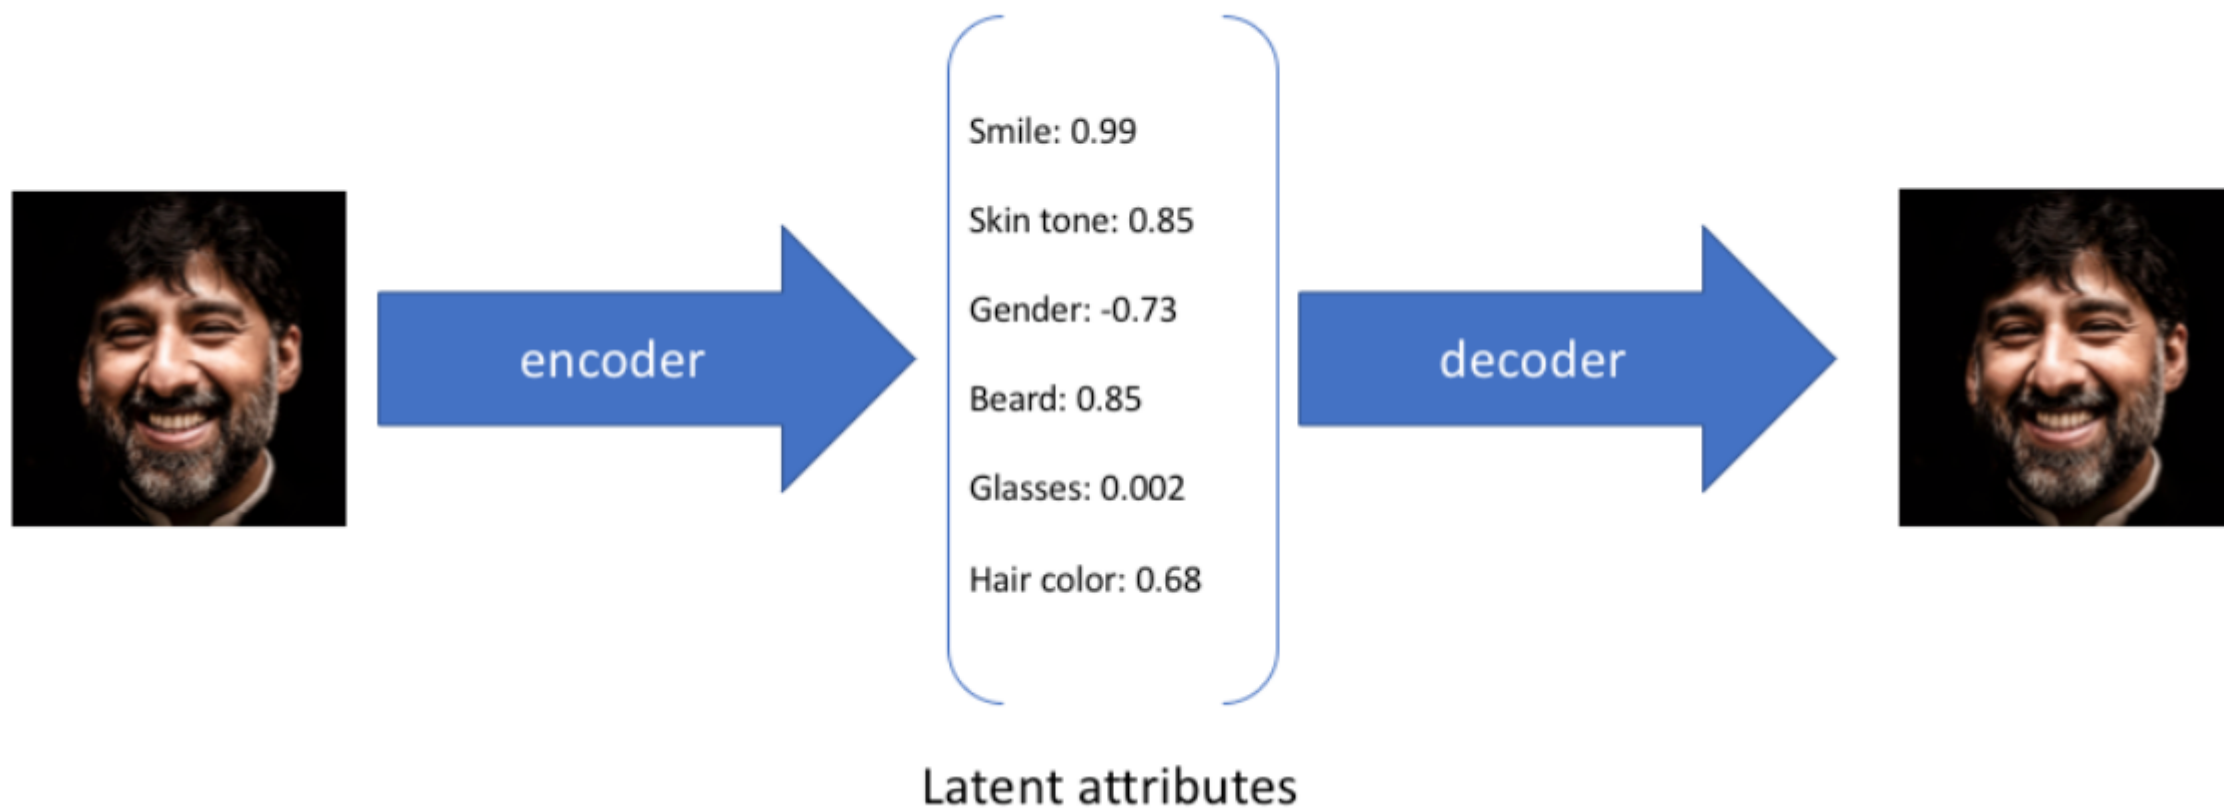
- Model is not able to develop a mapping which memorizes training data because input and target output are no longer same
- Rather, model learns a vector field for mapping input data towards a lower-dimensional manifold
  - If this manifold accurately describes the natural data, we have effectively "canceled out" added noise

# Variational Autoencoders (VAE)


- Proposed in 2013 by Diederik P. Kingma and Max Welling
- Provides a probabilistic manner for describing an observation in latent space
  - Trained to learn a mapping from input data to a probability distribution in a lower-dimensional latent space
  - Generate new samples from this distribution
  - Commonly used in image and text generation tasks
- Latent space is continuous
  - Decoder can generate new data points that interpolate among training data points
- Difference from autoencoder:
  - Provides a statistical manner for describing samples of dataset in latent space
  - Encoder outputs a probability distribution in bottleneck layer instead of a single output value




# AE





# VAE





# AE vs. VAE




Smile (discrete value)




Smile (probability distribution)



vs.



# AE vs VAE



# VAE

